

Sysdig

Container Security Workshop

CfgMgmt-2019

Miguel Julián @miky_kun

Julio García @papajulio

<https://github.com/figarocorso/security-workshop>



<Spoiler Alert>

- **Intro: who are we and what are we doing here?**
 - QA @ Sysdig
- **Container security best practices**
- **Kubernetes platform security features**
 - Run a few cases of use
- **Image scanning**
 - Scan a couple of images
 - Jenkins integration live demo
- **Runtime security and forensics**
- **Deploying the full set of open source tools**
 - Full demo with Sysdig Falco integrated with Nats + Kubeless



Few years ago...

First approach to docker

- It was told to be easy... until something goes wrong
 - How do I ssh there?
 - It's dead, so I cannot docker logs...

I'm an analog man, trapped in a digital world
Just an old school son of <beep> in this digital hell
Don't need no wifi, just want my hi-fi
Computers crashing, I want to smash 'em
I'm an analog man, dying in a digital world

Accept - Analog Man



Nowadays...

No stop signs

Speed limit

Nobody's gonna slow me down

Like a wheel

Gonna spin it

Nobody's gonna mess me around

Hey, Satan

Payin' my dues

Playin' in a rockin' band

Hey, mamma

Look at me

I'm on the way to the promised land

I'm on the highway to hell

AC/DC - Highway to Hell



Container security best practices



Resource usage

- Isolated does not mean they cannot disturb others
 - Our family has grown bigger with containers (vs. virtualization)
 - Thus we have to share better resources
 - DoS because of bug or malware
 - Many resources to share
 - CPU, RAM, storage, I/O, ...
 - But also file descriptors, user IDs, directory entries, ..
 - <https://sysdig.com/blog/container-isolation-gone-wrong/>
- `docker run -it --memory=2G --memory-swap=3G ubuntu bash`



Dockerfile

- **COPY >> ADD**
- **Secrets never in Dockerfile**
 - Nor at env variables, better use external tools like vaultproject
- **Do not forward privileged ports (22!)**
- **Force user directive, avoid using root user**
- **Avoid unnecessary packages**
 - Also clear as much as you can (keep images light)



Privileges

- **Enforce mandatory access control to prevent non desire access**
 - **seccomp, selinux, apparmor**
- **Create a non-root user**
 - **Remove setuid and setgui permission**
- **Do not --users=host (better isolate container users)**
- **--security-opt=no-new-privileges**
- **Make sure no aufs is used (buggy and not supported)**
- **Docker bench for security audit tool**
 - **<https://github.com/docker/docker-bench-security>**



Other considerations

- **Always pin image version**
 - Also, what about cached images? (`--no-cache` flag)
- **Image and container sprawl**
- **Building context is important (what are we mounting?)**
 - We can make use of `.dockerignore` file
 - Be careful not mounting `docker.sock` by mistake
- **Reduce attack surface (coreOS, RancherOS, Red Hat Atomic, ...)**
- **Avoid `docker0` in production (ARP spoofing, MAC flooding)**
- **`--pids-limit=100`, `--default-ulimit` and `--max-retries=5`**



Kubernetes



RBAC

- **Actors**
 - **Users**
 - **ServiceAccounts**
 - **Groups (SA prefix, user cert organization field)**
- **Resources**
 - **Pod, deployment, ...**
- **Role and ClusterRole**
 - **GET; WATCH; LIST; CREATE; UPDATE; PATCH; DELETE**
- **RoleBinding and ClusterRoleBinding**



K8S Security Policy

- **Pod/Container security context**
 - runAsUser, dropping capabilities, ...
- **Admission Controller**
 - NodeRestriction
 - ValidatingAdmissionWebhooks (integrate with Anchore)
- **Pod Security Policy**
 - Which kind of Pod can we create?
- **Kubernetes Network Policy**



About secrets and certs

- **Rotate, rotate, rotate**
 - Probably we might want to use an external service
 - Audit them ;-)
- **K8S only checks origin and expiration date**
- **How do I revoke access?**
 - Create a new user and remove privileges to old one
 - Recreate the CA and re-issue again the certs
- **What about ServiceAccounts tokens?**
 - Delete and create again



Let's have some fun

<https://github.com/figarocorso/security-workshop>

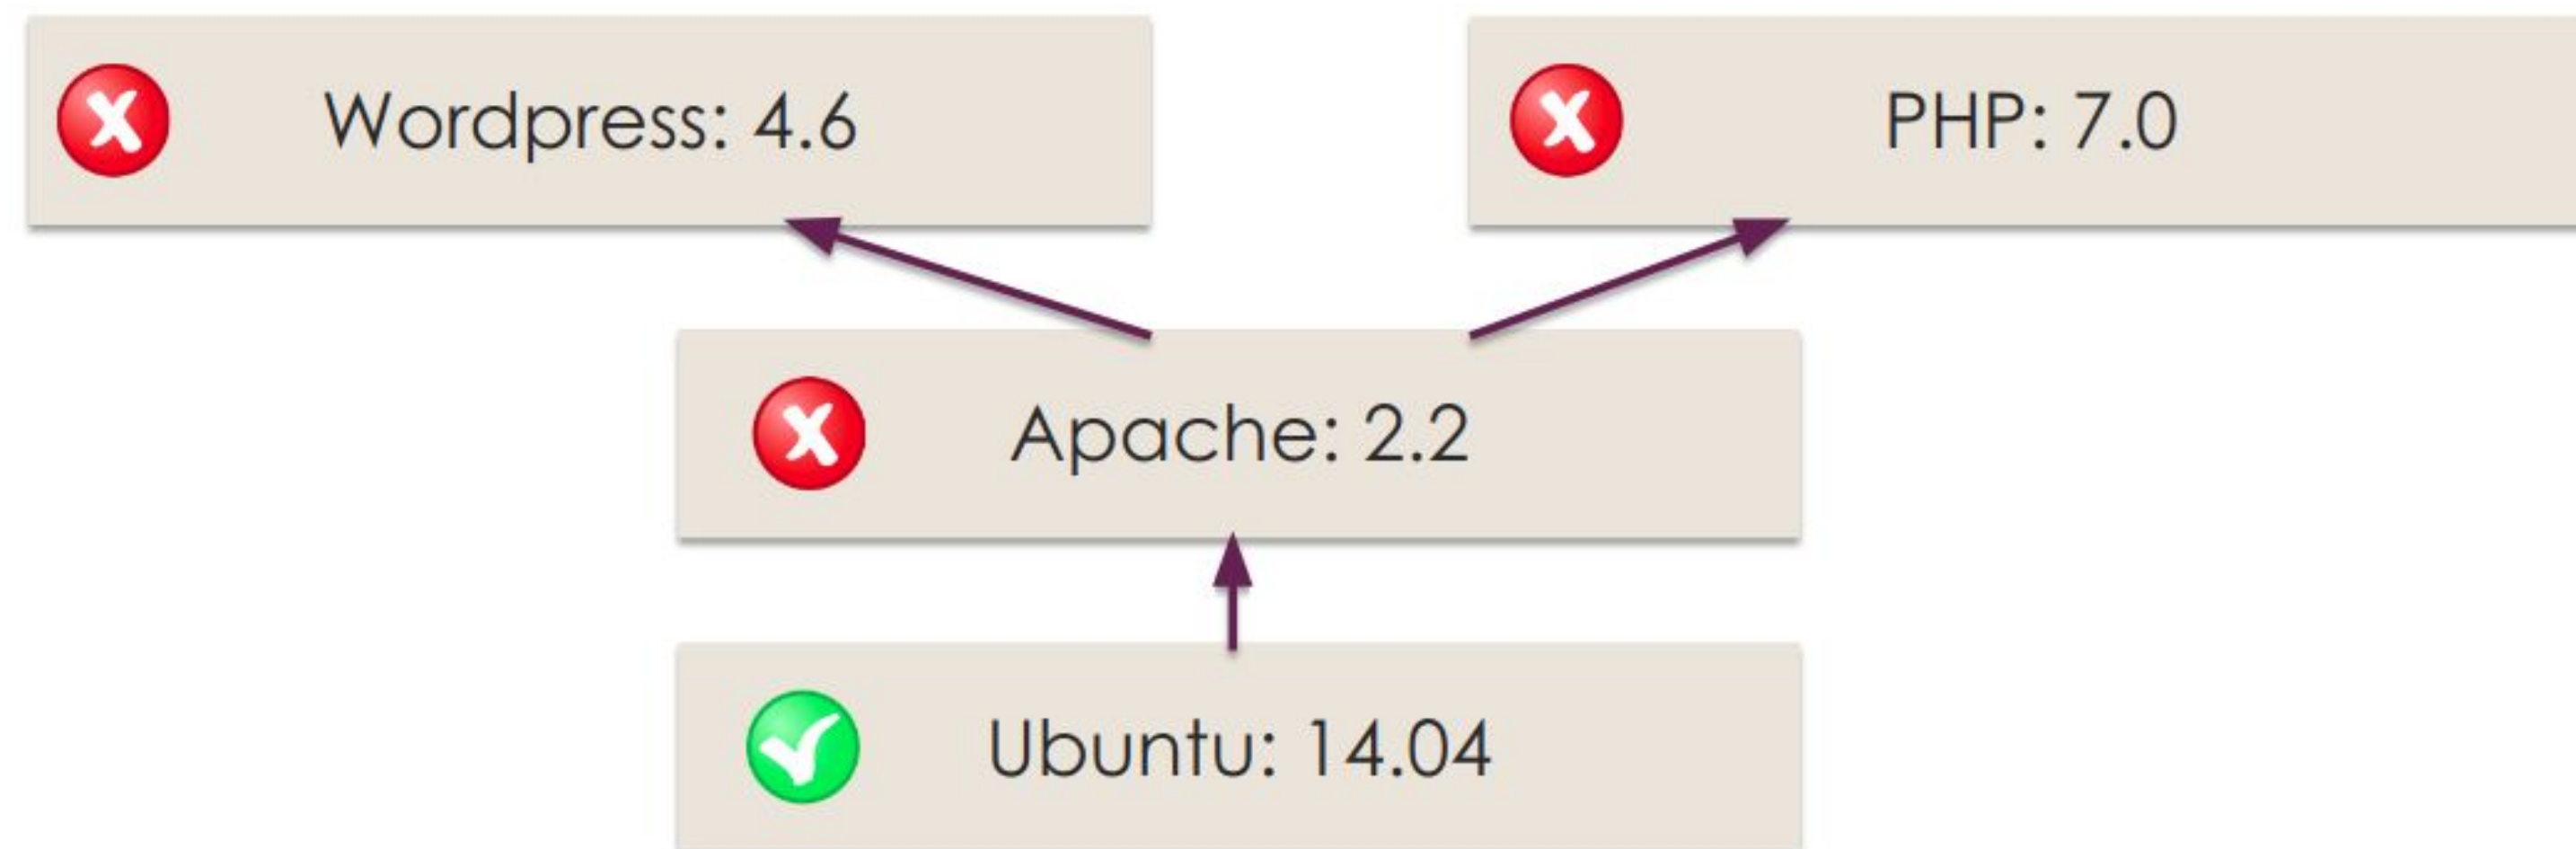


Image Scanning



Why image scanning?

- Define at our company what is secure to be run in our arch
- Layered design is quite cool but



Alternatives

- **Anchore Engine**
 - Centralized service for inspection, analysis and user defined rules
- **CoreOS/Clair**
 - Static analysis in application containers
- **Vuls.io**
 - Linux vulnerability scanner (info from NVD and OVAL)
 - Some container image support, but not container specific tool
- **OpenScap**
 - Audit tool following the NIST-certified SCAP
 - Again, not container specific

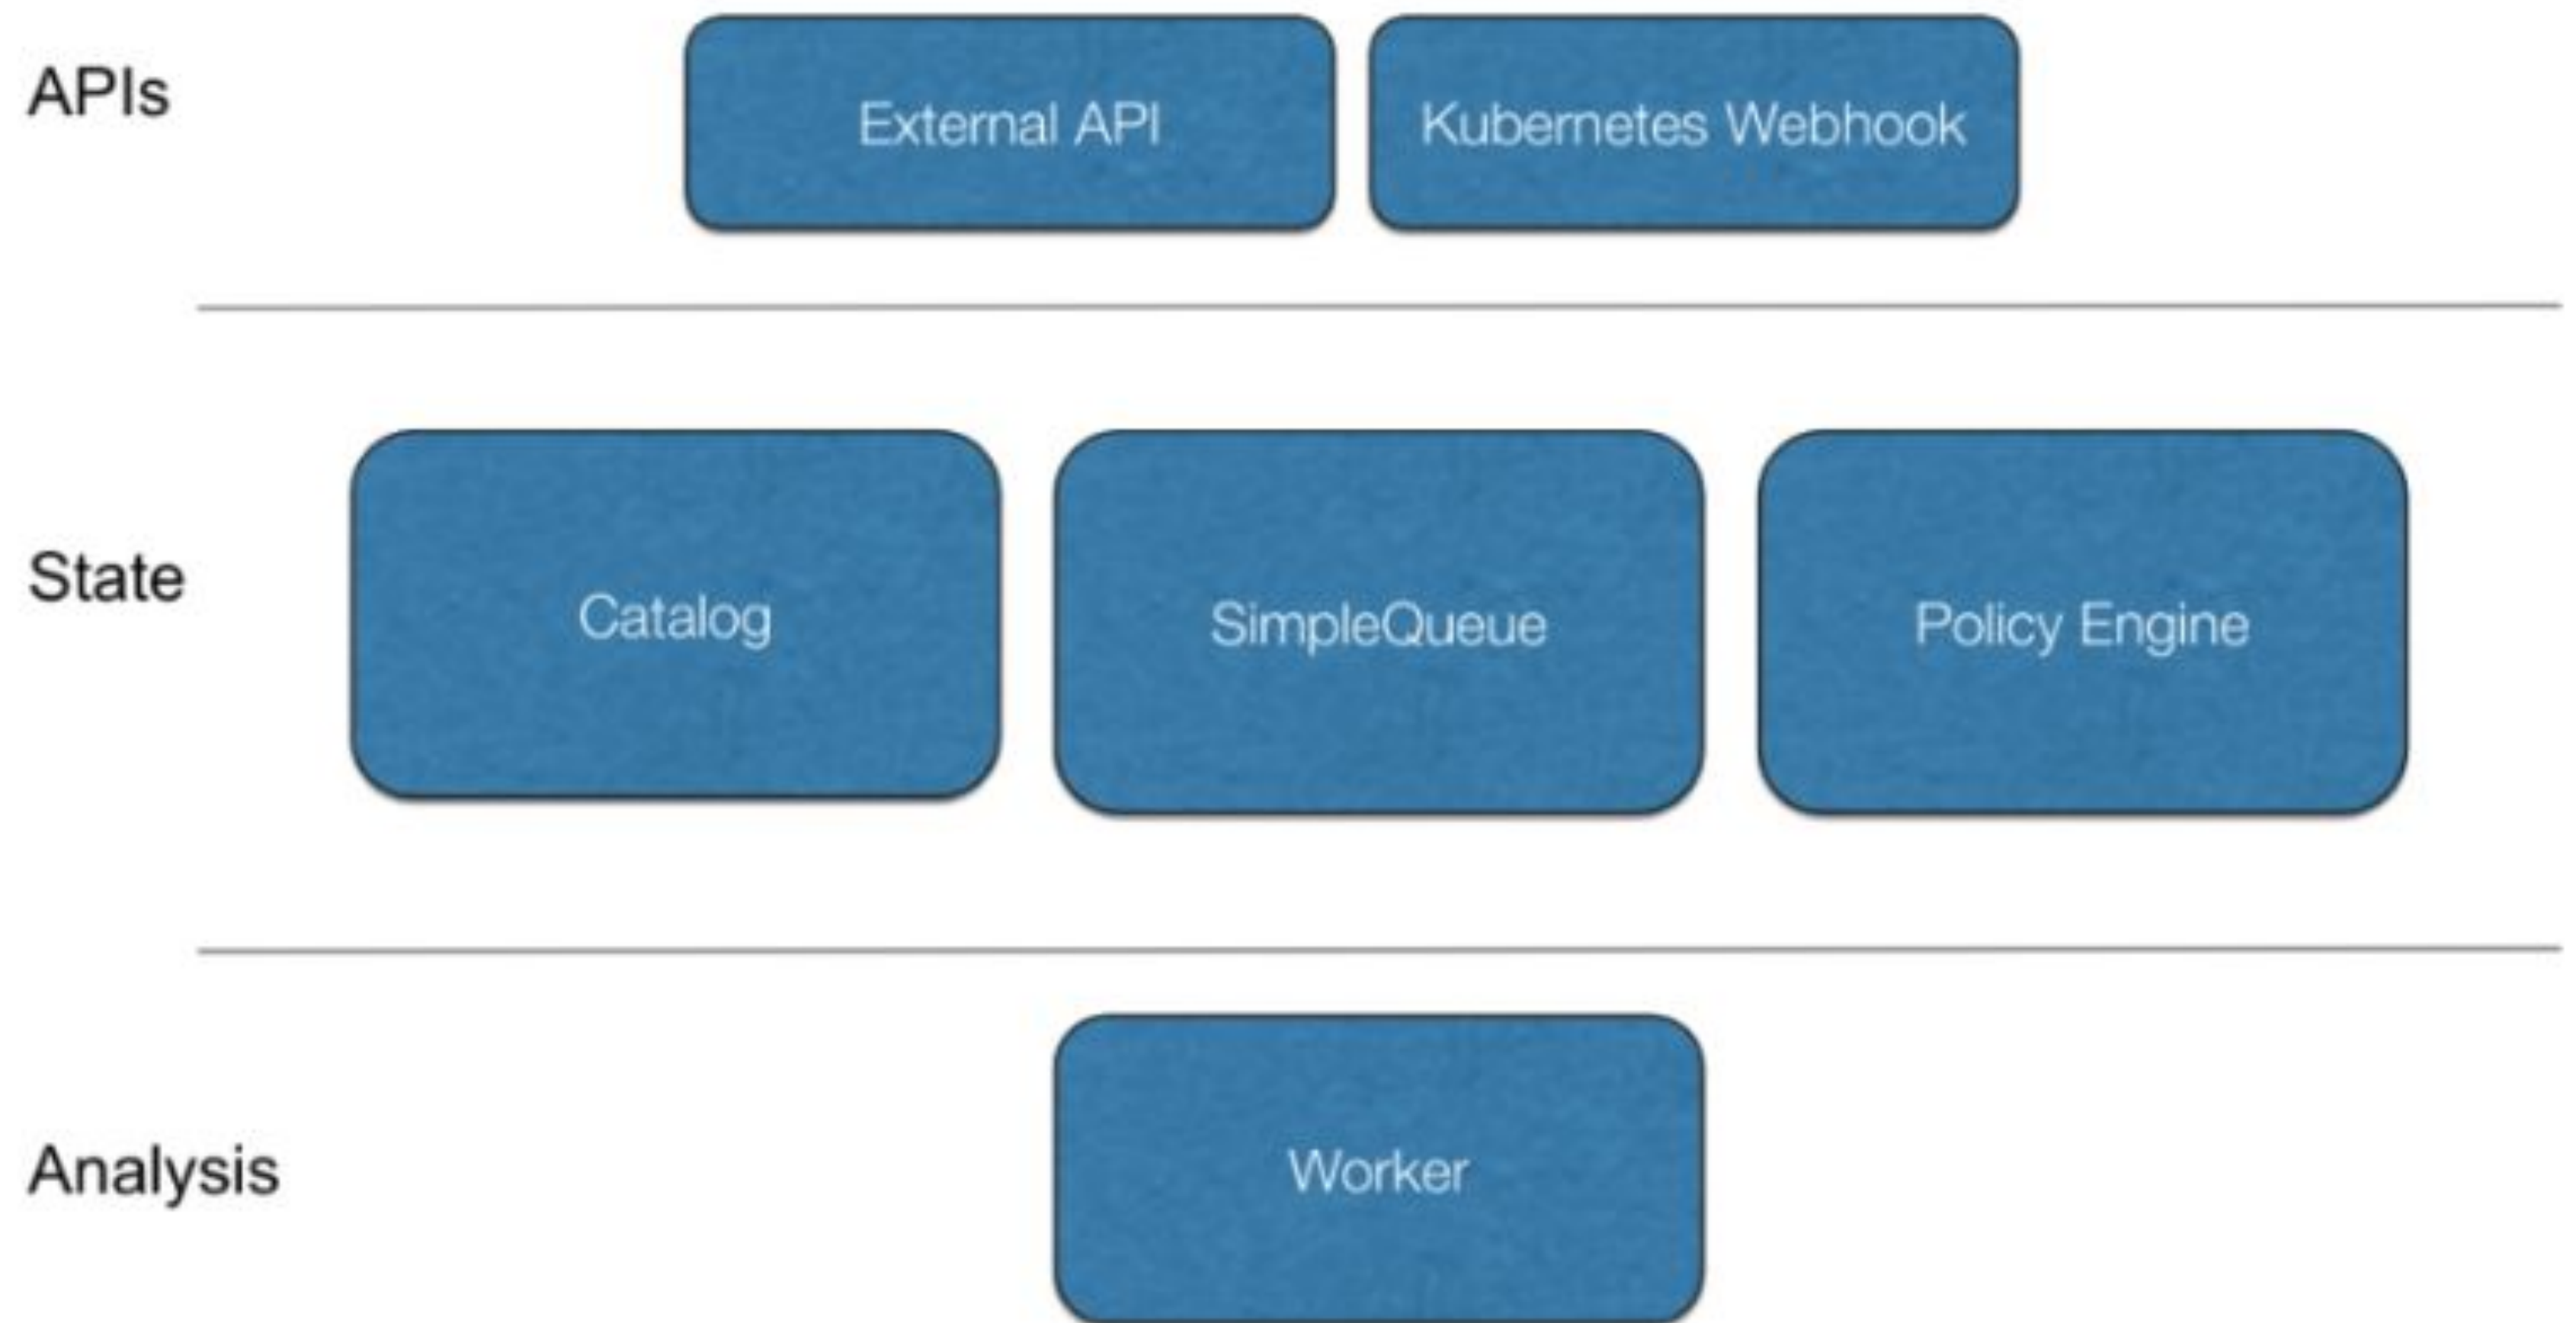


Image scanning

- Check pkgs, pypi, rubygem, binaries, OS files, ... against multiple vulnerabilities DB
- Whitelisted ones cached to speed-up process
- Analyze Dockerfile
 - Exposed ports
 - Privileged user (root)
 - Latest vs. pinned version
- User defined whitelists, blacklist, policies



Anchore architecture




Let's have some fun:

Deploy anchore



Let's do some Jenkins magic

—



Runtime security and forensics



Container security


- **Sandboxing: Seccomp**
 - Seccomp: strict mode (read/write to already opened files)
 - Seccomp bpf: policies (letting each process a few privileges)
- **Mandatory Access Control Systems: AppArmor**
 - Permitted linux capabilities
 - Permitted network operations
 - Allowed/Disallowed files
- **SELinux more complex, based in actor**

```
5 /usr/sbin/tcpdump {
6   #include <abstractions/base>
7   #include <abstractions/nameservice>
8   #include <abstractions/user-tmp>
9
10  capability net_raw,
11  capability setuid,
12  capability setgid,
13  capability dac_override,
14  network raw,
15  network packet,
16
17  # for -D
18  capability sys_module,
19  @{PROC}/bus/usb/ r,
20  @{PROC}/bus/usb/** r,
21
22  # for -F and -w
23  audit deny @{HOME}/.* mrwkl,
24  audit deny @{HOME}/.* / rw,
25  audit deny @{HOME}/.* /** mrwkl,
26  audit deny @{HOME}/bin/ rw,
27  audit deny @{HOME}/bin/** mrwkl,
28  @{HOME}/ r,
29  @{HOME}/** rw,
30
31  /usr/sbin/tcpdump r,
```



CONTAINER FORENSICS and INCIDENT RESPONSE


CONTAINER FORENSIC is DIFFICULT.
CONTAINERS CAN BE DESTROYED
RESCHEDULED into A DIFFERENT NODE



SECURITY INCIDENT

UNDERSTAND & CONTAIN
the IMPACT OF THE
SECURITY BREACH


- ✓ WHAT HAPPENED?
- ✓ WHAT WAS the BREACH?
- ✓ ANY DATA EXFILTRATED?
- ✓ HOW DID THEY BREAK the SYSTEM?
- ✓ WHO DID IT?



LIKE A
TIME MACHINE


COOL FEATURE!!

YOU NEED A WAY to PERFORM
FORENSICS.
EVEN IF YOUR CONTAINERS
ARE LONG GONE.



MEANINGFUL
TIMELINE

RECONSTRUCTING



REPORT

- EXPLOITED VULNERABILITIES
- SQL or COMMAND INJECTIONS
- WEAK CREDENTIALS
- INSECURE CONFIGURATIONS
- MALWARE & BACKDOORS
- CRYPTOMINING

INCIDENT
RESPONSE



1. CONTAIN the PROBLEM
2. MINIMIZE AFFECTED SERVICES
3. ELIMINATE the CAUSE of the PROBLEM

Behavioral monitoring

- **Falco!**
 - Look at system calls as a event stream
 - As a userspace process: context!
 - Easy to write rules
 - So we only filter and analyze that syscall event stream
 - This rulz, but we might be biased :-)

```
1 - rule: raw_network_socket
2   desc: an attempt to open a raw network socket by an unexpected program
3   condition: evt.type=socket and evt.dir=> and evt.arg.domain=AF_PACKET and not proc.name=tcpdump
4   output: Raw network socket opened by unexpected program (user=%user.name command=%proc.cmdline domain=%evt.arg.domain)
5   priority: WARNING
```



Reacting after an event

APPLICATION DEPLOYMENTS



KUBERNETES NODES

kubelet API



kubernetes



Kubeless

EXECUTE REACTION
i.e. kill the offending pod

F(x) F(x) F(x)


SUBSCRIBE TO
1..N TOPICS

WEBHOOK
NOTIFICATION

splunk[®]
phantom

EVENTS

K8S
METADATA



FALCO DAEMONSET



PUBLISH TO TOPIC

Let's have some fun:
Launch the whole scenario



One step forward:

Let's audit K8S events



Thank you!!

- Miguel Julián
 - @miky_kun
 - miguel.julian@sysdig.com
- Julio García
 - @papajulio
 - julio.garcia@sysdig.com

